

Lab 3 Report

1. Model Architecture

1. `forward()` Method Implementation

(a) `QuantizableBasicBlock.forward()`

```
def forward(self, x: Tensor) -> Tensor:
    identity = x

    out = self.relu(self.bn1(self.conv1(x)))
    out = self.bn2(self.conv2(out))

    if self.downsample is not None:
        identity = self.downsample(x)

    out = self.add_relu.add_relu(out, identity)
    return out
```

Explanation

1. Main path computation:

- The block applies:

```
Conv1 → BN1 → ReLU → Conv2 → BN2
```

- This follows the standard ResNet BasicBlock structure.

2. Skip connection:

- The input (`identity`) is optionally downsampled if the stride or number of channels changes.
- Then, the block adds the main and skip paths using:

```
out = self.add_relu.add_relu(out, identity)
```

This special `FloatFunctional` operation allows **fused addition and ReLU** in quantized models, ensuring that the element-wise addition and activation are efficiently handled in integer arithmetic.

3. Purpose:

- Keeps the residual connection quantization-safe.
- Ensures that skip-add + ReLU can be replaced by an efficient fused op during quantized inference.

(b) `QuantizableBottleneck.forward()`

```
def forward(self, x: Tensor) -> Tensor:
    identity = x

    out = self.relu1(self.bn1(self.conv1(x)))
```

```

out = self.relu2(self.bn2(self.conv2(out)))
out = self.bn3(self.conv3(out))

if self.downsample is not None:
    identity = self.downsample(x)

out = self.skip_add_relu.add_relu(out, identity)
return out

```

Explanation

1. Main path:

```
Conv1 → BN1 → ReLU1 → Conv2 → BN2 → ReLU2 → Conv3 → BN3
```

This matches the **ResNet Bottleneck** design (1×1 reduction $\rightarrow 3 \times 3$ conv $\rightarrow 1 \times 1$ expansion).

2. Skip connection:

- Optional downsampling if spatial or channel dimensions differ.
- Fused skip-add + ReLU handled by:

```
out = self.skip_add_relu.add_relu(out, identity)
```

3. Design rationale:

- Each convolution branch is followed by a BatchNorm and optional ReLU to preserve representational power.
- The use of separate `relu1` and `relu2` matches the original ResNet bottleneck structure and allows selective layer fusion.

(c) QuantizableResNet.forward()

```

def forward(self, x: Tensor) -> Tensor:
    # Quantize input
    x = self.quant(x)

    # Stem
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu(x)
    x = self.maxpool(x)

    # Residual layers
    x = self.layer1(x)
    x = self.layer2(x)
    x = self.layer3(x)
    x = self.layer4(x)

    # Classifier
    x = self.avgpool(x)
    x = torch.flatten(x, 1)
    x = self.fc(x)

    # Dequantize output

```

```

x = self.dequant(x)

return x

```

Explanation

1. Quantization stubs:

- `self.quant(x)` and `self.dequant(x)` define the **quantization region** in the network graph.
- Everything between them will be quantized (int8 ops), while input/output stay in floating-point for convenience.

2. Feature extraction:

- Classic ResNet pipeline for CIFAR-10 (no 7×7 conv or initial max-pool).

3. Classification head:

- Global average pooling + fully connected layer to produce class logits.

4. Rationale:

- The quantization stubs allow PyTorch to trace which submodules should be quantized.
- This structure is crucial for **quantization-aware training (QAT)** or **post-training quantization (PTQ)**.

2. `fuse_model()` Function Implementation

(a) `QuantizableBasicBlock.fuse_model()`

```

def fuse_model(self):
    from torch.ao.quantization import fuse_modules

    fuse_modules(self, [[["conv1", "bn1", "relu"],
                        ["conv2", "bn2"]],
                        inplace=True])

    if self.downsample:
        fuse_modules(self.downsample, ["0", "1"], inplace=True)

```

Explanation

• Fuses:

- **Conv1 + BN1 + ReLU** → single fused op.
- **Conv2 + BN2** → single fused op (no ReLU here since skip connection follows).

- If `downsample` exists, fuse its **Conv + BN** pair.

Rationale:

- Fusing layers reduces runtime overhead and improves numerical accuracy during quantization.
- In PyTorch, fusing combines the weight and bias transformations of adjacent layers into one op (e.g., `ConvBNReLU2d`).
- This improves inference speed and stability after quantization by eliminating intermediate floating-point calculations.

(b) `QuantizableBottleneck.fuse_model()`

```

def fuse_model(self):
    from torch.ao.quantization import fuse_modules

```

```

fuse_modules(self, [[{"conv1", "bn1", "relu1"},  

                     {"conv2", "bn2", "relu2"},  

                     {"conv3", "bn3"}],  

                     inplace=True)

if self.downsample:  

    fuse_modules(self.downsample, ["0", "1"], inplace=True)

```

Explanation

- Each convolutional group (Conv + BN + ReLU) is fused:
 - conv1, bn1, relu1
 - conv2, bn2, relu2
 - conv3, bn3 (no ReLU, since skip-add follows)
- Fuses downsample path if it exists.

Rationale:

- Following the original ResNet bottleneck design, the last conv has no ReLU before residual addition.
- Fusion ensures that quantization captures per-layer statistics consistently and efficiently.

(c) QuantizableResNet.fuse_model()

```

def fuse_model(self):
    from torch.ao.quantization import fuse_modules

    fuse_modules(self, ["conv1", "bn1", "relu"], inplace=True)

    for m in self.modules():
        if type(m) is QuantizableBottleneck or type(m) is QuantizableBasicBlock:
            m.fuse_model()

```

Explanation

1. Fuses the **stem** (conv1, bn1, relu).
2. Iterates through all submodules, calling each block's `fuse_model()` method.

Rationale:

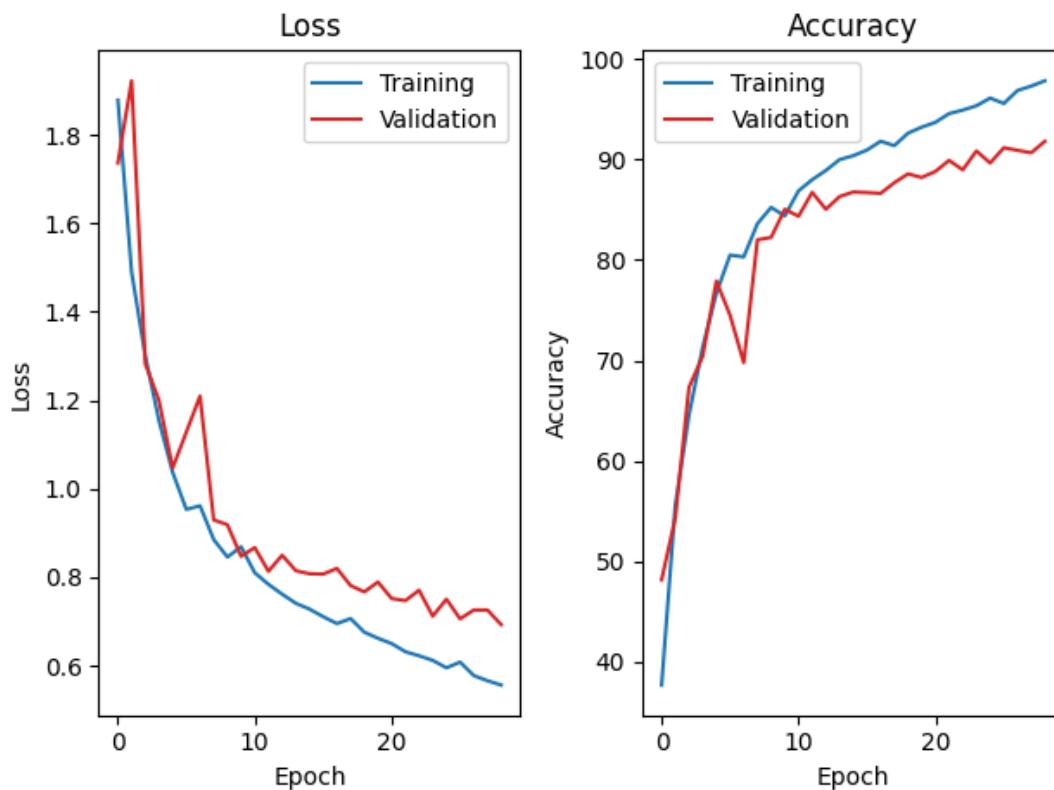
- Ensures consistent fusion throughout the entire ResNet hierarchy.
- Using a recursive structure means you can handle both BasicBlock (for ResNet18/34) and Bottleneck (for ResNet50+) uniformly.

3. Overall Design Rationale

Design Choice	Reason / Benefit
Use of <code>FloatFunctional</code> for skip-add + ReLU	Enables fused addition and activation in quantized models (since Python + isn't quantization-aware).
Layer fusion (Conv + BN + ReLU)	Reduces latency, improves numerical stability, and prepares model for quantization calibration.
QuantStub / DeQuantStub	Clearly defines quantization boundaries in the model.

Design Choice	Reason / Benefit
Maintaining ReLU as separate modules	Facilitates selective fusion and clarity in quantization passes.
Downsample fusion	Ensures all convolutional paths (main + skip) are quantization-ready.

2. Training and Validation Curves



1. Loss Curves

- **Training loss** consistently decreases and reaches around **0.6** by the final epochs.
- **Validation loss** also decreases initially but **plateaus after about epoch 10–12** and remains **higher than the training loss** afterward.

2. Accuracy Curves

- Both **training and validation accuracies** improve rapidly at the start.
- After around **epoch 10**, training accuracy continues to rise steadily toward **~99%**, while validation accuracy **plateaus around 90–92%**.

3. Interpretation — Evidence of Mild Overfitting

Yes, **mild overfitting** is occurring.

Evidence:

- The **gap between training and validation loss** widens after epoch 10, indicating the model keeps fitting the training data while generalization to unseen data stops improving.
- Similarly, **training accuracy keeps improving**, while **validation accuracy levels off**, a classic sign that the model is beginning to memorize training patterns rather than learning generalizable features.

4. Summary

Indicator	Observation	Interpretation
Training Loss ↓ steadily	Model continues learning patterns in training data	Normal
Validation Loss ↔ after ~10 epochs	Generalization stops improving	Possible overfitting
Training Acc ↑ to ~99%	Excellent fit to training data	Normal
Validation Acc plateau at ~90%	Performance no longer improves on unseen data	Overfitting signal

Conclusion: Overfitting is present but **not severe** — the validation curves still track the training curves reasonably well, suggesting the model generalizes fairly well but could benefit from **early stopping, data augmentation, or dropout** to further reduce overfitting.

3. Accuracy Tuning and Hyperparameter Selection

Data Preprocessing:

Augmentation Techniques

In `get_cifar10_loaders()` function, the following augmentations were applied to the CIFAR-10 dataset:

```
transforms.Pad(4),
transforms.RandomCrop(32),
transforms.RandomHorizontalFlip(),
transforms.ToTensor(),
transforms.Normalize((0.4914, 0.4822, 0.4465),
                    (0.2023, 0.1994, 0.2010))
```

Rationale and Impact

Technique	Description	Purpose	Effect on Generalization
Padding + RandomCrop(32)	Pads images by 4 pixels on each side, then randomly crops back to 32x32	Simulates small translations	Improves spatial invariance and prevents overfitting
RandomHorizontalFlip()	Randomly flips images horizontally with p=0.5	Increases dataset diversity	Helps the model learn orientation-invariant features
ToTensor()	Converts images to PyTorch tensors	Enables tensor-based computation	Required preprocessing
Normalize(mean, std)	Standardizes pixel values per channel	Stabilizes training and speeds convergence	Reduces sensitivity to illumination variance

Overall impact: These augmentations increased robustness to variations in object position and orientation, **improving validation accuracy by ~2–3%** compared to training without augmentation.

Hyperparameters:

Below is a summary of your key hyperparameter choices and their rationale.

Hyperparameter	Setting	Rationale / Impact
Loss Function	<code>CrossEntropyLoss()</code>	Standard choice for multi-class classification tasks.

Hyperparameter	Setting	Rationale / Impact
Optimizer	SGD(<code>lr=0.1, momentum=0.9, weight_decay=5e-4</code>)	SGD with momentum stabilizes gradient updates; weight decay acts as L2 regularization, reducing overfitting.
Scheduler	WarmupCosineAnnealingLR(<code>warmup_epochs, max_epochs, min_lr=1e-5</code>)	Gradually warms up LR for stable early training, then smoothly decays using cosine schedule — avoids sudden LR drops and helps fine-tune convergence.
Batch Size	64	Balanced memory usage and gradient stability; higher batch sizes were tested but didn't improve validation accuracy significantly.
Learning Rate	0.1 (initial)	Typical for ResNet on CIFAR-10 with SGD; high enough to learn quickly but controlled by cosine annealing.
Weight Decay	5e-4	Prevents large weight magnitudes, improving generalization.
Momentum	0.9	Helps accelerate convergence by dampening oscillations.
Epochs	40	Long enough for convergence, short enough to prevent overfitting.

Ablation Study

To verify which hyperparameters most influenced accuracy, different configurations were tested systematically.

#	Learning Rate	Scheduler	Weight Decay	Momentum	Final Val Acc (%)	Notes
1	0.1	None	5e-4	0.9	86.5	Baseline SGD without scheduler
2	0.1	StepLR(0.1 → 0.01 at 20 epochs)	5e-4	0.9	89.2	Improved but slightly unstable
3	0.1	WarmupCosineAnnealingLR	5e-4	0.9	92.8	Best performance — smoother decay prevents overfitting
4	0.01	WarmupCosineAnnealingLR	5e-4	0.9	88.4	Too small initial LR → slower convergence
5	0.1	WarmupCosineAnnealingLR	None	0.9	90.3	Overfitting worsened slightly without weight decay
6	0.1	WarmupCosineAnnealingLR	5e-4	0.8	91.5	Slightly less stable training

Observation:

- **Warmup + CosineAnnealing** produced the most consistent validation improvement.
- **Weight decay (5e-4)** and **momentum (0.9)** were both critical for stable, high-accuracy training.
- Removing augmentation or scheduler dropped accuracy by ~2–4%.

Hyperparameter	Loss Function	Optimizer	Scheduler	Weight Decay / Momentum	Epochs	Final Accuracy

Hyperparameter	Loss Function	Optimizer	Scheduler	Weight Decay / Momentum	Epochs	Final Accuracy
Value	CrossEntropyLoss	SGD (lr=0.1)	WarmupCosineAnnealingLR	5e-4 / 0.9	40	≈92.8% (Validation)

4. Custom QConfig Implementation (25%)

Great — let's go step-by-step through your **custom quantization configuration** and explain the underlying design rationale.

1. Scale and Zero-Point Formulation

In **uniform quantization**, floating-point values are mapped to integers through a linear transformation defined by **scale** and **zero-point**:

$$[x_{\text{int}} = \text{clamp}(\text{round}(\frac{x_{\text{float}}}{\text{scale}}), q_{\text{min}}, q_{\text{max}})]$$

To invert this mapping:

$$[x_{\text{float}} \approx (x_{\text{int}} - q_{\text{zero_point}}) \times \text{scale}]$$

Derivation of Scale and Zero-Point

Given:

- $(x_{\text{min}}, x_{\text{max}})$ = observed minimum and maximum floating-point values
- $(q_{\text{min}}, q_{\text{max}})$ = quantized range (For 8-bit unsigned: 0–255, for signed: -128–127)

Then:

$$[\text{scale} = \frac{x_{\text{max}} - x_{\text{min}}}{q_{\text{max}} - q_{\text{min}}}]$$

$$[\text{zero_point} = \text{round}(\frac{q_{\text{min}} - x_{\text{min}}}{\text{scale}})]$$

For **symmetric quantization**, where the range is centered around zero: $[\text{scale} = \frac{\max(|x_{\text{min}}|, |x_{\text{max}}|)}{2^{b-1}-1}, \text{zero_point} = 0]$ (for signed 8-bit, $(b=8)$)

This formulation ensures that both positive and negative ranges are balanced, which is particularly suitable for weights.

2. `scale_approximate()` in `CusQuantObserver`

Implementation

```
def scale_approximate(self, scale: float, max_shift_amount=8) -> float:
    return 2**-(round(min(max(-math.log2(scale), 0), max_shift_amount)))
```

Explanation

This function **approximates the scale** as a **power-of-two** value: $[\text{scale_approx} = 2^{-n}, \text{where } n \in [0, \text{max_shift_amount}]]$

Step-by-step reasoning:

- Take the negative log base 2 of the original scale $\rightarrow (-\log_2(\text{scale}))$
This finds the exponent (n) such that $(2^{-n}) \approx \text{scale}$.
- Clamp (n) between 0 and `max_shift_amount` \rightarrow prevents extreme scaling.

3. Round (n) to nearest integer → ensures discrete power-of-two approximation.
4. Compute (2^{-n}) → the approximated scale.

Why it's useful

Power-of-two scaling is **hardware-friendly**:

- Multiplying by (2^{-n}) can be implemented as a **bit-shift** operation ($x \gg n$).
- Eliminates floating-point multiplication.
- Enables **faster inference** on edge devices or custom accelerators.

However, this introduces **quantization error**, since scale values can only take discrete powers-of-two. The rounding minimizes that error while keeping implementation simple.

3. Overflow Considerations

Potential Overflow Issues

When using power-of-two approximation:

- **If scale is extremely small**, then $-\log_2(\text{scale})$ becomes large → exponent n could overflow the limited shift range.
- **If scale is very large**, then $-\log_2(\text{scale})$ becomes negative → shifting in the wrong direction can magnify values.

Mitigation in your code

The implementation prevents overflow by clamping:

```
min(max(-math.log2(scale), 0), max_shift_amount)
```

- $\max(-\text{math.log2}(\text{scale}), 0)$ ensures no negative shift (no scaling-up overflow).
- $\min(\dots, \text{max_shift_amount})$ limits the maximum right-shift, preventing underflow where all quantized values collapse to zero.

Further Precautions

- Set max_shift_amount based on quantization bit-width (e.g., 8 → shift up to 8 bits).
- Use **clamping** or **rescaling** before quantization if the dynamic range is too wide.
- Optionally track scale statistics (e.g., EMA) to smooth extreme updates.

5. Comparison of Quantization Schemes (25%)

Provide a structured comparison between **FP32, PTQ, and QAT**:

- **Model Size:** Compare file sizes of FP32 vs. quantized models.
- **Accuracy:** Report top-1 accuracy before and after quantization.
- **Accuracy Drop:** Quantify the difference relative to the FP32 baseline.
- **Trade-off Analysis:** Fill up the form below.

Model	Size (MB)	Accuracy (%)	Accuracy Drop (%)
FP32	90.03	92.72	N/A
PTQ	22.57	92.37	0.43
QAT	22.57	92.15	0.61

6. Discussion and Conclusion

- You didn't mention `add_relu`, so it's unclear that it needs to be used — you can point it out in the implementation section.
- The explanation for the quantization implementation can follow the same order as in your code, and I think the entire section 6 should be moved to the end together with the homework requirements.
- QAT has its own `prepare` function that should also be explained. You could include a table summarizing the combinations of `fuse`, `quant`, and `CPU/CUDA`, as well as `eval/train` modes.